

Problem Set 11

Your name and student ID

Today's date

Run this chunk of code to load the autograder package!

Instructions

- Solutions will be released by Sunday, November 24th for exam studying.
- This semester, homework assignments are for practice only and will not be turned in for marks.

Helpful hints:

- Every function you need to use was taught during lecture! So you may need to revisit the lecture code to help you along by opening the relevant files on Datahub. Alternatively, you may wish to view the code in the condensed PDFs posted on the course website. Good luck!
- Knit your file early and often to minimize knitting errors! If you copy and paste code for the slides, you are bound to get an error that is hard to diagnose. Typing out the code is the way to smooth knitting! We recommend knitting your file each time after you write a few sentences/add a new code chunk, so you can detect the source of knitting errors more easily. This will save you and the GSIs from frustration!
- If your code runs off the page of the knitted PDF then you will LOSE POINTS! To avoid this, have a look at your knitted PDF and ensure all the code fits in the file (you can easily view it on Gradescope via the provided link after submitting). If it doesn't look right, go back to your .Rmd file and add spaces (new lines) using the return or enter key so that the code runs onto the next line.

Parental leave is often compensated to some degree, but the amount of compensation varies greatly. You read a research article that stated, “across people of all incomes, 47% of leave-takers received full pay during their leave, 16% received partial pay, and 37% received no pay.”

After reading this, you wonder what the distribution of parental leave payment is for low income households. Suppose you conduct a survey of leave-takers within households earning less than \$30,000 per year. You surveyed 225 people (selected via a random sample) and found that 51 received full pay, 33 received partial pay, and 141 received no pay.

1. You would like to investigate whether the distribution of pay for households earning < \$30,000 is different from that of all income levels. Does this correspond to a chi-square test of independence or a chi-square test for goodness of fit?

This corresponds to a chi-square test for goodness of fit. This is because we only have one categorical variable (a sample from low income households) and are comparing the observed counts for each category to a provided distribution.

2. What are the expected counts of leave-takers among households with incomes < \$30,000? Assign each expected count to the appropriate variable and to 2 decimal places.

```
full_pay <- 105.75
partial_pay <- 36.00
no_pay <- 83.25
```

```
. = ottr::check("tests/p2.R")
```

```
##
```

```
## All tests passed!
```

3. State the null hypothesis in the context of the question.

H_0 : The distribution of leave-takers is the same as that in the research article (i.e., the proportion receiving full pay equals 47%, the proportion receiving partial pay is 16%, and the proportion with no pay is 37%).

4. Compute the chi-square statistic. Round your answer to 2 decimal places.

```
chi_sq <- 68.66  
chi_sq
```

```
## [1] 68.66
```

```
. = ottr::check("tests/p4.R")
```

```
##
```

```
## All tests passed!
```

$$\chi^2 = \sum_{i=1}^k \frac{(O_i - E_i)^2}{E_i}$$

$$= (105.75 - 51)^2/105.75 + (36 - 33)^2/36 + (83.25 - 141)^2/83.25 = 28.34574 + 0.25 + 40.06081 = 68.65656$$

5. Uncomment the cell (i.e., the term in the summation) that contributes the most to the test statistic.

```
# UNCOMMENT THE CORRECT ANSWER

# largest_contribution <- 'full pay'
# largest_contribution <- 'partial pay'

largest_contribution <- 'no pay'
```

The largest contribution comes from the deviation in the people that receive no pay to go on parental leave. We see a much higher number of no pay among low income households than that expected under the null hypothesis.

```
. = ottr::check("tests/p5.R")
```

```
##
```

```
## All tests passed!
```

6. Compute the p-value. Round your answer to 2 decimal places.

```
p_value <- round(pchisq(q = 68.65656, df = 2, lower.tail = F), 2)
p_value <- 0.00
p_value
```

```
## [1] 0
```

```
. = ottr::check("tests/p6.R")
```

```
##
```

```
## All tests passed!
```

7. Is there evidence against the null hypothesis? Uncomment your selection below.

```
# UNCOMMENT THE CORRECT ANSWER
```

```
# conclusion <- 'in favor of null'
```

```
conclusion <- 'against null'
```

The probability of seeing this chi-square statistic is very tiny (<0.001) under the null hypothesis. Thus we conclude there is evidence in favor of the alternative hypothesis that the distribution of leave is different for low income households vs. that specified in the research article.

```
. = ottr::check("tests/p7.R")
```

```
##
```

```
## All tests passed!
```


Human papillomavirus (HPV) is a very common STI that most sexually active persons will encounter during their lifetimes. While many people clear the virus, certain strands can lead to adverse health outcomes such as genital warts and cervical cancer.

Suppose that you selected a random sample from a population and collected these data on age and HPV status for the sample:

Age Group	HPV +	HPV -	Row total
14-19	160	492	652 (33.9%)
20-24	85	104	189 (9.8%)
25-29	48	126	174 (9.1%)
30-39	90	238	328 (17.1%)
40-49	82	242	324 (16.9%)
50-59	50	204	254 (13.2%)
Col total	515 (26.8%)	1406 (73.2%)	1921

8. What are the explanatory and response variables? Uncomment the appropriate answer.

```
# UNCOMMENT THE CORRECT ANSWER
```

```
variable_type <- c('explanatory: age group', 'response: HPV status')
```

```
# variable_type <- c('explanatory: HPV status', 'response: age group')
```

```
. = ottr::check("tests/p8.R")
```

```
##
```

```
## All tests passed!
```


9. State the null and alternative hypotheses to test whether there is a relationship between age group and HPV status.

H_0 : The conditional distribution of age is the same for HPV + and HPV - individuals.

H_A : The conditional distribution of age is different for HPV + and HPV - individuals.

10. Run the code below to examine the conditional distribution of age by HPV status. Based on this plot, which age group will contribute the most to the chi-square statistic? Explain why.

(That is, can you tell based on this plot when the observed count will differ most from the expected count under the null hypothesis of no relationship between age group and HPV status?)



Cells corresponding to the 20-24 year-olds will likely contribute the most to the chi-square statistic because they exhibit the largest observed difference between HPV- and HPV+ individuals. (Additionally, one might mention that the low overall proportion for 20-24 year olds means the denominator for the 20-24 year old chi-square term will be relatively small).

11. Fill out the table of expected counts under the null hypothesis of no association between age group and HPV status. Round each number to 2 decimal places.

Expected counts:

Age Group	HPV +	HPV -
14-19	$652 \cdot 515 / 1921 = 174.7944$	$652 \cdot 1406 / 1921 = 477.2056$
20-24	$189 \cdot 515 / 1921 = 50.66892$	$189 \cdot 1406 / 1921 = 138.3311$
25-29	$174 \cdot 515 / 1921 = 46.64758$	$174 \cdot 1406 / 1921 = 127.3524$
30-39	$328 \cdot 515 / 1921 = 87.93337$	$328 \cdot 1406 / 1921 = 240.0666$
40-49	$324 \cdot 515 / 1921 = 86.86101$	$324 \cdot 1406 / 1921 = 237.139$
50-59	$254 \cdot 515 / 1921 = 68.09474$	$254 \cdot 1406 / 1921 = 185.9053$

```
A <- 174.79
B <- 50.67
C <- 46.65
D <- 87.93
E <- 86.86
G <- 68.09
H <- 477.21
I <- 138.33
J <- 127.35
K <- 240.07
L <- 237.14
M <- 185.91
```

```
. = ottr::check("tests/p11.R")
```

```
##
## All tests passed!
```

12. Calculate the test statistic. Round your answer to 2 decimal places.

```
chi_sq_p12 <- 40.55  
chi_sq_p12
```

```
## [1] 40.55
```

$$\chi^2 = \sum_{i=1}^k \frac{(O_i - E_i)^2}{E_i}$$

$$\begin{aligned} &= [(174.7944 - 160)^2/174.7944] + [(477.2056 - 492)^2/477.2056] + [(50.66892 - 85)^2/50.66892] + \\ &[(138.3311 - 104)^2/138.3311] + [(46.64758 - 48)^2/46.64758] + [(127.3524 - 126)^2/127.3524] + [(87.93337 - \\ &90)^2/87.93337] + [(240.0666 - 238)^2/240.0666] + [(86.86101 - 82)^2/86.86101] + [(237.139 - 242)^2/237.139] \\ &+ [(68.09474 - 50)^2/68.09474] + [(185.9053 - 204)^2/185.9053] = 1.252181 + 0.4586582 + 23.26126 + \\ &8.520314 + 0.03920975 + 0.01436161 + 0.04857041 + 0.01779021 + 0.2720371 + 0.09964334 + 4.808295 + \\ &1.761209 = 40.55353 \end{aligned}$$

```
. = ottr::check("tests/p12.R")
```

```
##
```

```
## All tests passed!
```

13. Calculate the p-value for your test statistic. Round your answer to 2 decimal places.

```
df <- (6-1)*(2-1) # 5
p_value_p13 <- round(pchisq(q = 40.55353, df = 5, lower.tail = F), 2)
p_value_p13 <- 0.00
p_value_p13
```

```
## [1] 0
```

```
. = ottr::check("tests/p13.R")
```

```
##
```

```
## All tests passed!
```

14. Is there evidence against the null hypothesis? Uncomment your selection.

```
# UNCOMMENT THE CORRECT ANSWER  
# conclusion_p14 <- 'in favor of null'  
  
conclusion_p14 <- 'against null'
```

END